

See discussions, stats, and author profiles for this publication at: <https://www.researchgate.net/publication/302532508>

# Cloud for Gaming

Chapter · May 2015

DOI: 10.1007/978-3-319-08234-9\_39-1

CITATIONS

0

READS

903

3 authors, including:



**Gabriele D'Angelo**  
University of Bologna

106 PUBLICATIONS 911 CITATIONS

[SEE PROFILE](#)



**Stefano Ferretti**  
University of Bologna

231 PUBLICATIONS 6,095 CITATIONS

[SEE PROFILE](#)

Some of the authors of this publication are also working on these related projects:



Number of cancer in Italy [View project](#)



E-learning Accessibility [View project](#)

# Cloud for Gaming

Gabriele D'Angelo      Stefano Ferretti      Moreno Marzolla  
 University of Bologna, Dept. of Computer Science and Engineering  
 Mura Anteo Zamboni 7, I-40126 Bologna, Italy  
 {g.dangelo, s.ferretti, moreno.marzolla}@unibo.it

## Definition

Cloud for Gaming refers to the use of cloud computing technologies to build large-scale gaming infrastructures, with the goal of improving scalability and responsiveness, improve the user's experience and enable new business models.

## What is cloud computing?

Cloud computing is a service model where the provider offers computation and storage resources to customers on a “pay as you go” basis [12]. The essential features of a cloud computing environment are:

- On-demand self service:** the ability to provide computing capabilities (e.g., CPU time, network storage) dynamically, as needed, without human intervention;
- Broad network access:** resources can be accessed through the network by client platforms using standard mechanisms and protocols;
- Resource pooling:** virtual and physical resources can be pooled and assigned dynamically to consumers, according to their demand, using a multi-tenant model;
- Elasticity:** from the customers point of view, the provider offers unlimited resources that can be purchased in any quantity at any time;
- Measured service:** cloud resource and service usages are optimized through a pay-per-use business model, and are monitored, controlled and reported transparently to both their customer and provider.

The typical interaction between cloud provider and customer works as follows: the customer connects to a “cloud marketplace” through a Web interface, and selects the type and amount of the resources she needs (e.g., some virtual servers with given number of CPU cores, memory and disk space). The resources are allocated from a large pool that is physically hosted on some big datacenter managed by the cloud provider. Once instantiated, the resources are accessed by the customer through the network. Additional resources can be acquired at a later time, e.g., to cope with an increase of the workload, and released when no longer needed. The customer pays a price that depends on the type and amount of resources requested (e.g., CPU cores speed, memory size, disk space), and on the duration of their usage.

The *service model* defines the level of abstraction at which the cloud infrastructure provides service (Figure 1). In a *Software as a Service* (SaaS) cloud, the system provides application services running in the cloud. “Google apps” is an example of a widely used SaaS cloud. In contrast, the capabilities provided by a *Platform as a Service* (PaaS) cloud consist of programming languages, tools and a hosting environment for applications developed by the customer. The difference




Figure 1: Cloud Service Model




Figure 2: Gaming as a Service models

between the SaaS and PaaS models is that while the user of a SaaS cloud simply utilizes an application that runs in the cloud, the user of a PaaS cloud develops an application that can be executed in the cloud and made available to service customers; the application development is carried out using libraries, APIs and tools possibly offered by some other company. Examples of PaaS solutions are AppEngine by Google, Force.com from Salesforce, Microsoft's Azure and Amazon's Elastic Beanstalk. Finally, an *Infrastructure as a Service* (IaaS) cloud provides its customers with fundamental computing capabilities such as processing, storage and networks where the customer can run arbitrary software, including operating systems and applications. The number of companies offering such kind of services is continually growing; one of the earliest being Amazon with its EC2 platform.

The *deployment model* defines the mode of operation of a cloud infrastructure; these are the private cloud, the community cloud, the public cloud, and the hybrid cloud models. A *private cloud* is operated exclusively for a customer organization; it is not necessarily managed by that organization. In the *community cloud* model the infrastructure is shared by several organizations and supports a specific community with common concerns (e.g., security requirements, policy enforcement). In the *public cloud* model the infrastructure is made available to the general public and is owned by an organization selling cloud services. Finally, the *hybrid cloud* model refers to cloud infrastructures constructed out of two or more private, public or community clouds.

## Cloud computing for gaming

The gaming industry embraced the cloud computing paradigm by implementing the *Gaming as a Service* (GaaS) model [2]. Different instances of the GaaS paradigm have been proposed: remote rendering GaaS, local rendering GaaS and cognitive resource allocation GaaS.

In the *remote rendering* GaaS (RR-GaaS) model the cloud infrastructure hosts one instance of the game engine for each player (Fig. 2 (a)). An encoder module running on the cloud is responsible for rendering every frame of the game scene, and compressing the video stream so that it can be transmitted to the user's terminal where the stream is decoded and displayed. User inputs are acquired from the terminal and sent back to the game engine that takes care of updating the game




Figure 3: Number of online players of the Runescape MMOG; the data refers to the period from may 5 to may 16, 2011

state accordingly. The advantage of the RR-GaaS model is that the workload on the terminal is greatly reduced, since the computationally demanding step of rendering the game scenes is entirely offloaded to the cloud. This allows complex games to be played on less powerful devices, such as mobile phones or cheap game consoles, that are only required to be capable of decoding the video stream in real time. However, the RR-GaaS model consumes considerable bandwidth to transmit the compressed video stream, and may be particularly sensitive to network delays. Examples of RR-GaaS implementations are GamingAnywhere [8] and Nvidia Grid<sup>TM</sup><sup>1</sup>.

In the *local rendering* GaaS model, the video stream is encoded on the cloud as a sequence of high-level rendering instructions that are streamed to the player terminal (Fig. 2 (b)); the terminal decodes and executes the instructions to draw each frame. Since encoding of each frame as a sequence of drawing instructions is often more space efficient than compressing the resulting bitmap, the LR-GaaS model may require less network bandwidth than RR-GaaS, and therefore eliminate the need for real-time video transmission capability. This comes at the cost of requiring a more powerful terminal with an adequate graphics subsystem.

Finally, in the *cognitive resource allocation* GaaS model, the game engine is logically partitioned into a set of modules that can be upload and executed at the client side (Fig. 2 (c)). As the game evolves, the terminal receives and executes the appropriate modules, and may keep or discard the unused ones. The CRA-GaaS model shifts the computation back to the client terminal, therefore reducing the load on the cloud. However, the client resources are used efficiently, since at any time only the needed components are stored locally. This is a significant advantage if we consider that the data of a complete modern game takes a lot of space for textures, 3D models, sounds and code modules.

GaaS provides advantages for both game developers and players. The ability to offload some computation on the cloud allows simple terminals such as mobile devices to play complex games. Since the game engine is accessed on demand, flexible business models such as pay-per-play or monthly subscription can be easily implemented. Finally, game operators can scale up and down the amount of cloud resources used by the gaming infrastructure.

The last point is particularly important, especially for the so-called *Massively Multiplayer On-*

<sup>1</sup><http://www.nvidia.com/object/cloud-gaming.html>, Accessed on 2015/4/4

*line Games* (MMOG). Modern MMOGs are large-scale distributed systems serving millions of concurrent users which interact in real-time with a large, dynamic virtual world. The number of users playing the game at any given time follows a pattern that originates from the typical daily human activity. As an example, Figure 3 shows the number of online players of RuneScape<sup>2</sup> [10], a fantasy game where players can travel across a fictional medieval realm. During the observed period, more than 200,000 players are connected to the system at peak hours; this number reduces to about 110,000 players during off-peak hours. Hence, the daily churn (number of players leaving/joining the system during the day) is about 100,000 users. It is evident that static resource provisioning based on the average load results in system overloaded roughly half the time; provisioning for the worst case results in a massive resource under-utilization.

To effectively implement a cloud-based gaming infrastructure, it is necessary to address non-trivial issues related to game state partitioning, responsiveness, synchronization and security.

**Partitioning** The key factor for achieving scalability of a GaaS infrastructure is the ability to partition the workload across the cloud resources. This is relatively easy if the workload consists of the execution of independent game instances that can be executed on any available resource, irrespective of where other instances are running. This is the case when the game does not allow different players to interact. Things become complex if the instances are *not* independent, as in the case of a MMOG system where all players interact with the same virtual world. In this case, the game engine must maintain a large shared state, allowing the players to “see” the effects of actions performed by the other players operating in the same virtual location.

This is achieved by partitioning the virtual world across multiple zones, each handled by a separate set of cloud resources. Given that communication between resource instances may incur significant delays, it is important that interaction across neighboring zones is minimized. For example, each partition may hold a collection of “islands” such that all interactions happen within the collection, while players can jump from one “island” to another.

Depending on the (virtual) mobility pattern of each player, some areas of the game field may become crowded, while others may become less populated. In order to cope with this variability, each zone controller is physically hosted on resources provided and operated by a cloud infrastructure. The cloud provider is in general a separate entity providing computational and storage resources to the game operator on a pay-as-you go model. This means that the game operator can request additional servers and/or additional storage space at any time, and release them when no longer needed. Thus, the game operator can request more resources when the workload on a zone increases, in order to keep the response time perceived by players below a predefined maximum threshold. When the workload decreases, the game operator can release surplus resources in order to reduce costs.

**Synchronization** The success of a gaming system is based on having players perceiving the game state as identical and simultaneously evolving on every player participating to a gaming session. If the game state is replicated in different cloud servers, a synchronization algorithm is needed to maintain the consistency of the redundant game state. To this aim, different schemes have been proposed in the literature [1]. They mainly differ from classic synchronization algorithms employed by distributed systems in their additional requirement for keeping the computation quick and responsive. To this aim, some schemes relax the requirements for full consistency during the game state computation.

A basic distinction is between conservative and optimistic synchronization. Conservative synchronization approaches allow the processing of game updates only when it is consistency-safe to do so. Lockstep [6], time-bucket synchronization [6], interactivity restoring [5], are some examples in the literature.

Optimistic synchronization mechanisms process game updates as soon as they receive them, thus increasing the responsiveness of the system. Yet, it is assumed that most updates are received in the correct order and that, in any case, it would be acceptable to recover later from possible

---

<sup>2</sup><http://www.runescape.com>

inconsistencies. Examples of optimistic approaches available in the scientific literature are the optimistic bucket synchronization [4], the combination of local lag and Time Warp proposed in [11], the trailing state synchronization [3], and the improved Time Warp equipped with the dropping scheme and a correlation-based delivery control approach [5].

**Responsiveness** The task of providing a pleasant experience to players becomes challenging when trying to deploy a large scale and highly interactive online game. Responsiveness means having small delays between the generation of a game update at a given player and the time at which all other players perceive such update. How much such delays must be small depends on the type of online game. Obviously, the shorter the delay the better is. But it is possible to identify a game-specific responsiveness threshold  $T_r$  that represents the maximum delay allowable before providing a game update to players. The typical  $T_r$  for fast-paced games (e.g., first-person shooter, racing vehicles) is 150 to 200ms, but this value can be increased to seconds in slow paced-games (e.g., strategic, role-playing games) [5, 13].

A key point is that each player is geographically distributed. Thus, his latency to reach the game server on the cloud is usually different from other players. If a classic client-server approach is employed, it might thus happen that a responsive service is provided to some subset of users, while the other players can perceive a non responsive game evolution. This raises another main issue, i.e. fairness provision. This means guaranteeing that all players have the same chance of winning, regardless of their subjective network conditions [5]. To this aim, it should be guaranteed that all players perceive the same and simultaneous game evolution at the same time.

GaaS infrastructures represent an effective tool to provide responsive and fair gaming experiences. Cloud servers can manage the game state evolution in a scalable manner. Multiple server instances can be run in the same datacenter, when needed. Moreover, if the game involves world wide distributed players, one might think to introduce a federation of cloud servers, geographically distributed, so that each client/player might connect to its nearest server. This could balance the network delays between the player and its server, thus augmenting the fairness level provided by the system. However, when multiple servers are involved, each one with a redundant copy of the game state, synchronization algorithm are needed to maintain game state consistency.

**Security and reliability** The security issues of GaaS infrastructures have become mainstream after the PlayStation Network outage that, in 2011, has halted the Sony online gaming network for 23 days. The network was shut down after detecting an external intrusion that led to a huge number of accounts being compromised, and the exposure of the players' personal information.

From the reliability point of view, large cloud systems provide some level of redundancy to cope with failures, including the use of geographically distributed datacenters, so that catastrophic events do not cause a complete outage. Unfortunately, the GaaS infrastructure may still represent a single point of failure; the PlayStation Network outage is just one example: in that case a security incident prompted the system administrators to temporarily shut down the whole service. Other possibilities must be considered as well: for example, the company operating the GaaS infrastructure may go bankrupt, depriving all players from the game service they might already have paid for.

From the security point of view, GaaS infrastructures are affected by the typical issues of cloud computing (e.g. insider attacks [14]) and online gaming (e.g. cheating [7]). Online games are an appealing target for hacks because players often invest huge amount of time in their character development, and is therefore quite easy to monetize game items on the black market. Additionally, individual accounts on online gaming platforms often contains information, such as credit card numbers, that are the typical target of cyber-criminals. Details of the avatar of each player can provide information such as sexual preferences [9] that could cause considerable embarrassment if made public.

## References

- [1] Game event synchronization. In Borko Furht, editor, *Encyclopedia of Multimedia*, pages 256–257. Springer US, 2006.
- [2] Wei Cai, Min Chen, and Victor C. M. Leung. Toward gaming as a service. *IEEE Internet Computing*, 18(3):12–18, May 2014.
- [3] Eric Cronin, Burton Filstrup, Anthony R. Kurc, and Sugih Jamin. An efficient synchronization mechanism for mirrored game architectures. In *Proceedings of the 1st Workshop on Network and System Support for Games*, NetGames '02, pages 67–73, New York, NY, USA, 2002. ACM.
- [4] C. Diot and L. Gautier. A distributed architecture for multiplayer interactive applications on the internet. *Network, IEEE*, 13(4):6–15, Jul 1999.
- [5] Stefano Ferretti. *Synchronization in Multiplayer Online Games*, pages 175–196. John Wiley & Sons, Inc., 2014.
- [6] Richard M. Fujimoto. *Parallel and Distribution Simulation Systems*. John Wiley & Sons, Inc., New York, NY, USA, 1st edition, 1999.
- [7] Jiankun Hu and Fabio Zambetta. Security issues in massive online games. *Security and Communication Networks*, 1(1):83–92, 2008.
- [8] Chun-Ying Huang, Kuan-Ta Chen, De-Yu Chen, Hwai-Jung Hsu, and Cheng-Hsin Hsu. Gaminganywhere: The first open source cloud gaming system. *ACM Trans. Multimedia Comput. Commun. Appl.*, 10(1s):10:1–10:25, January 2014.
- [9] Searle Huh and Dmitri Williams. Dude looks like a lady: Gender swapping in an online game. In William Sims Bainbridge, editor, *Online Worlds: Convergence of the Real and the Virtual*, Human-Computer Interaction Series, pages 161–174. Springer London, 2010.
- [10] Moreno Marzolla, Stefano Ferretti, and Gabriele D’Angelo. Dynamic resource provisioning for cloud-based gaming infrastructures. *Computers in Entertainment*, 10(1):4:1–4:20, December 2012.
- [11] Martin Mauve, Stefan Fischer, and Jörg Widmer. A generic proxy system for networked computer games. In *Proceedings of the 1st Workshop on Network and System Support for Games*, NetGames '02, pages 25–28, New York, NY, USA, 2002. ACM.
- [12] P. Mell and T. Grance. The NIST Definition of Cloud Computing (Draft)–Recommendations of the National Institute of Standards and Technology. Special publication 800-145 (draft), Gaithersburg (MD), January 2011.
- [13] Lothar Pantel and Lars C. Wolf. On the impact of delay on real-time multiplayer games. In *Proceedings of the 12th International Workshop on Network and Operating Systems Support for Digital Audio and Video*, NOSSDAV '02, pages 23–29, New York, NY, USA, 2002. ACM.
- [14] Dimitrios Zissis and Dimitrios Lekkas. Addressing cloud computing security issues. *Future Generation Computer Systems*, 28(3):583 – 592, 2012.